

# Offensive and Defensive Cybersecurity

–

***Heap Exploitation***

21-22

# Heap Introduction

- What's the heap?
  - One or more memory pages used to store data(rw-)
- Why do we need it?
  - For dynamic memory allocation
  - What about alloca, mmap, etc? Ok, but ...
- How do we manage it?
  - Through library functions

# Memory Allocations

- **libc**

- malloc - allocate a chunk of memory
- calloc - allocate and zero-out memory
- realloc - change size of an allocation
- free - free a chunk of memory

- **syscall**

- mmap (allocate memory page)
- munmap (deallocate memory page)
- brk/sbrk (change the location of the program break)

# The HEAP Allocators

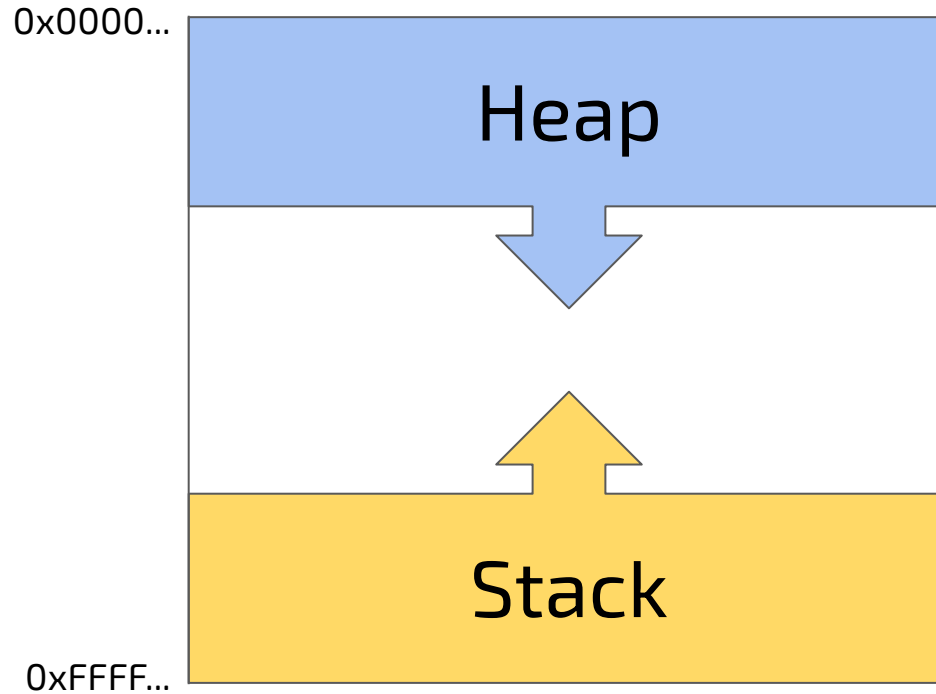
- **ptmalloc** (glibc)
- dlmalloc (was in glibc)
- tcmalloc (chromium)
- jemalloc (FreeBSD, Firefox, Android)

splittings, fits, coalescing, segregations (free list, storage, non determinism)

## ptmalloc2 (aka the malloc of glibc)

- **splittings** (how to divide in chunk)
- **fits** (match requested size with )
- **coalescing** (how to merge chunks)
- **segregations** free list
- NO segregations storage
- **deterministic**

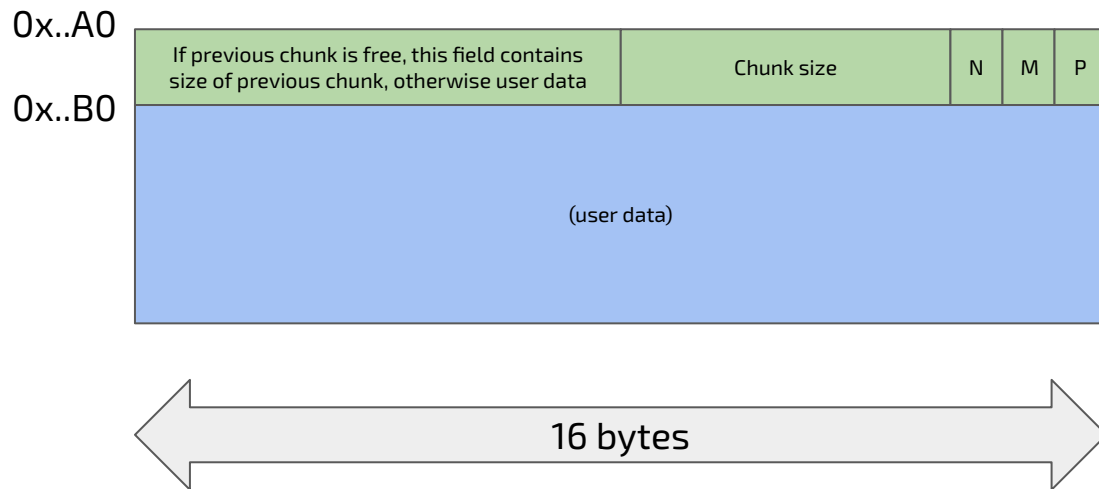
# Memory Layout



## Memory Allocation: malloc

- `void* malloc(size_t size);`
- Most known allocation primitive
- Requires allocation size(bytes)
- Returns a “void \*” pointer which points the allocated memory(buffer)
- This buffer is a part of struct called chunk

# Chunk



- PREV\_INUSE (P) – This bit is set when previous chunk is allocated.
- IS\_MMAPPED (M) – This bit is set when chunk is mmap'd.
- NON\_MAIN\_ARENA (N) – This bit is set when this chunk belongs to a thread arena.



## Top Chunk

- Special chunk that occupies all the available memory space in the heap
- Every time a malloc is called it might be shrunk
- Once there's no more space on the heap, a `brk(void *)` is called to allocate more pages to the heap and the top chunk is expanded

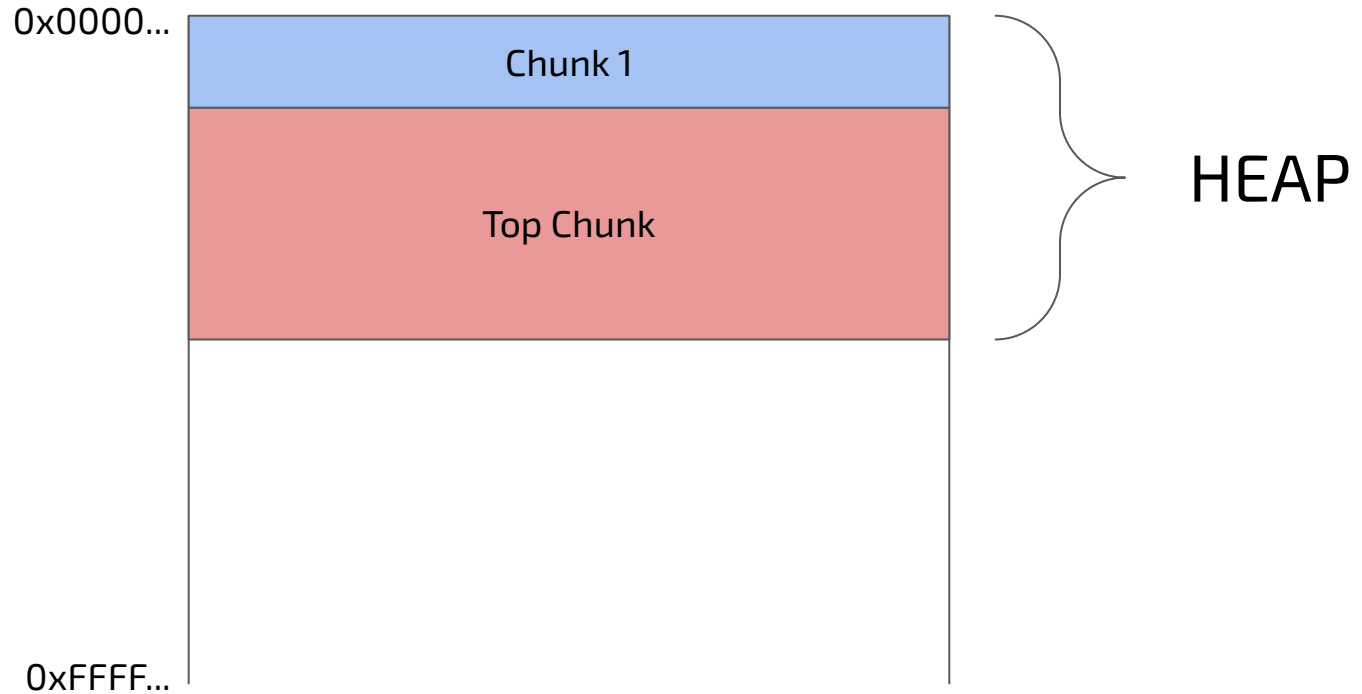
# Memory Allocation

0x0000...

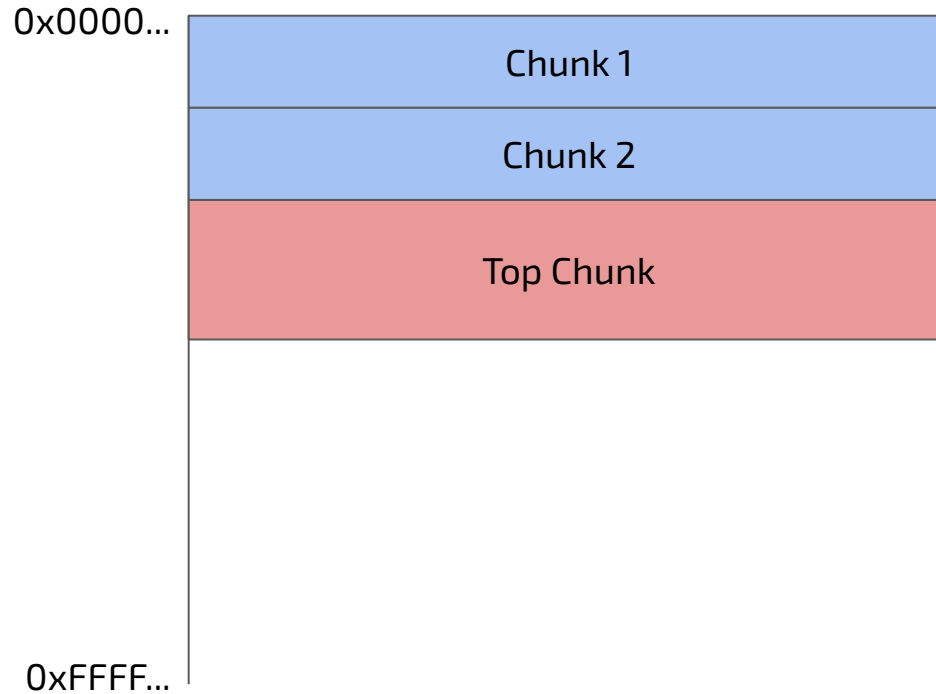
0xFFFF...



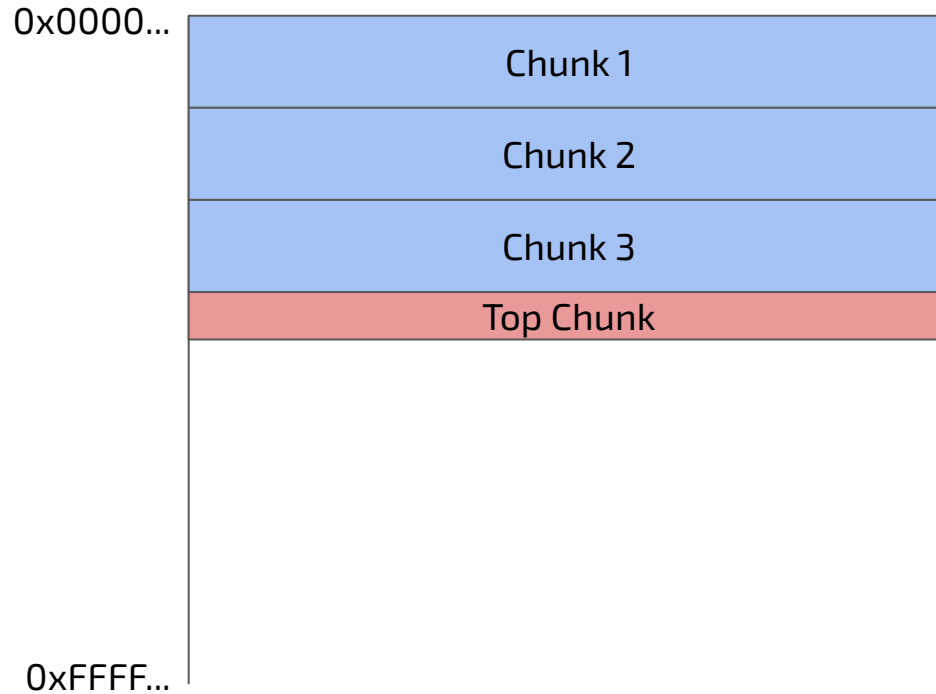
# Memory Allocation



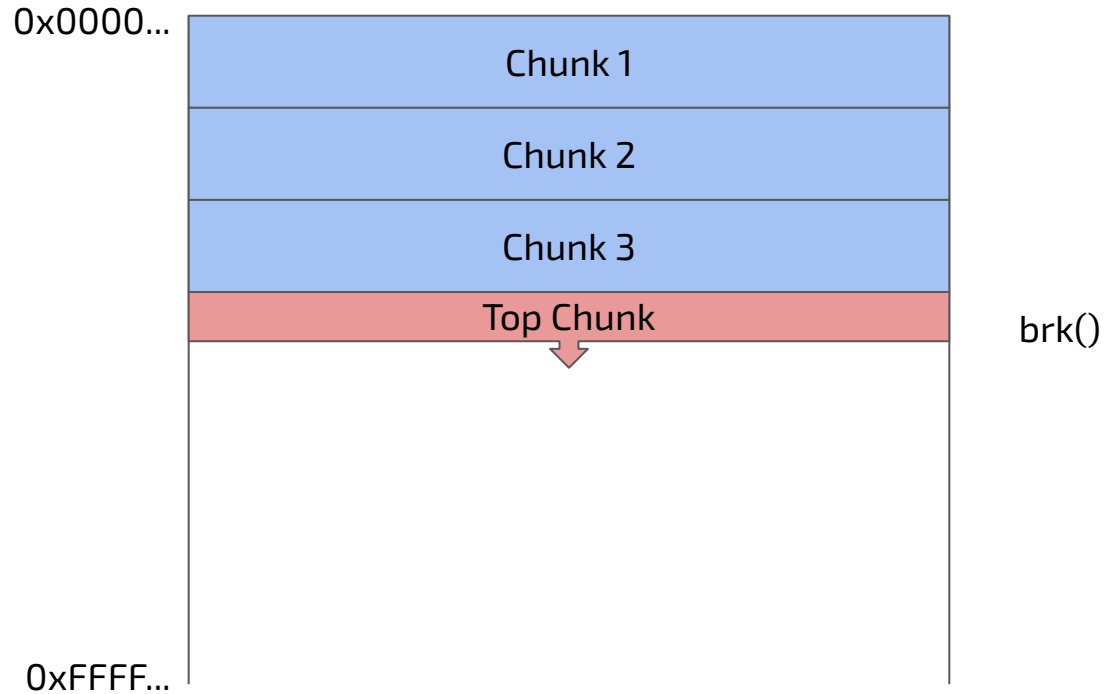
# Memory Allocation



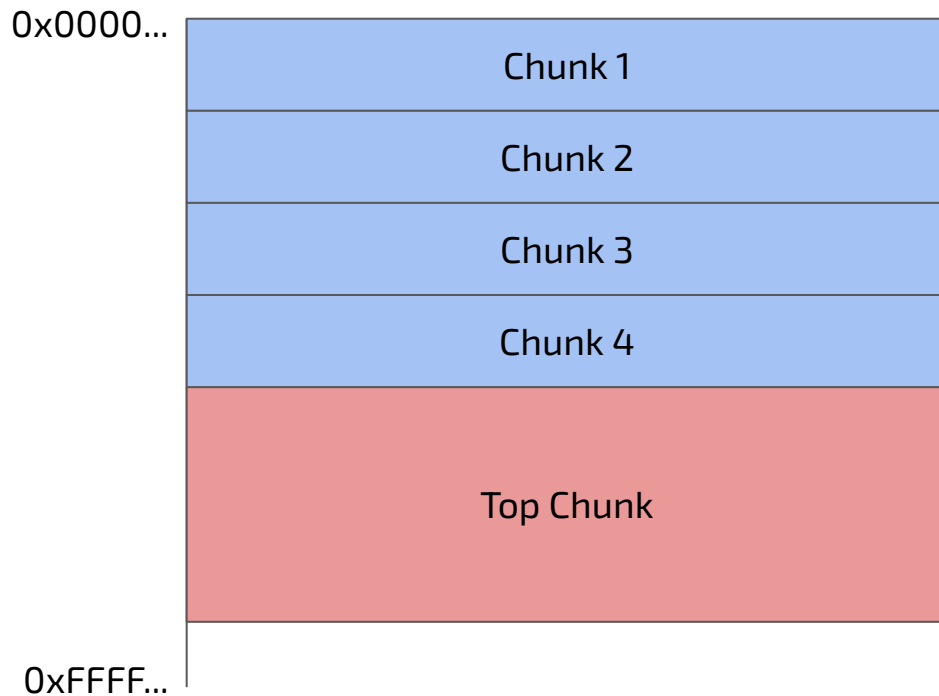
# Memory Allocation



# Memory Allocation



# Memory Allocation



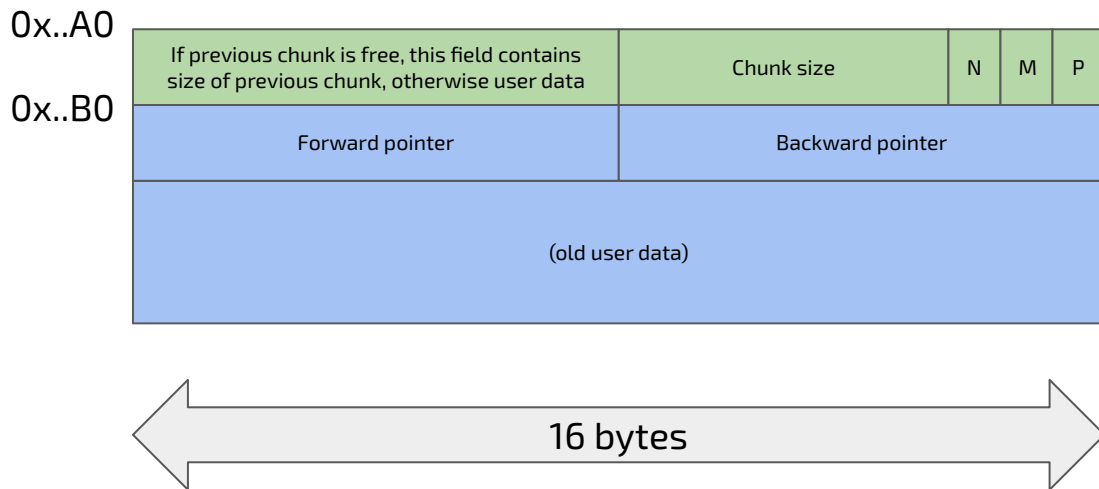
Let's see it in memory!



# Memory Deallocation: free

- `void free(void* ptr);`
- Most known deallocation primitive
- Requires a pointer to a memory buffer previously allocated with a function for memory allocation (e.g. `malloc`)
- Freed chunks could be consolidated with other freed chunks (also with the top chunks)
- If not they are inserted in lists called bins

# Free Chunk



- PREV\_INUSE (P) – This bit is set when previous chunk is allocated.
- IS\_MMAPPED (M) – This bit is set when chunk is mmap'd.
- NON\_MAIN\_ARENA (N) – This bit is set when this chunk belongs to a thread arena.

# Bins

- Lists of freed chunks of a specific size
- Heads of the lists are located in the .bss of the libc (main\_arena)
- Lists can be single or double linked
- 4 types of bins:
  - Fast bins – 8 Linked lists
  - Unsorted bin – 1 Double linked list
  - Small bins – 62 Double linked lists
  - Large bins – 62 Double linked lists

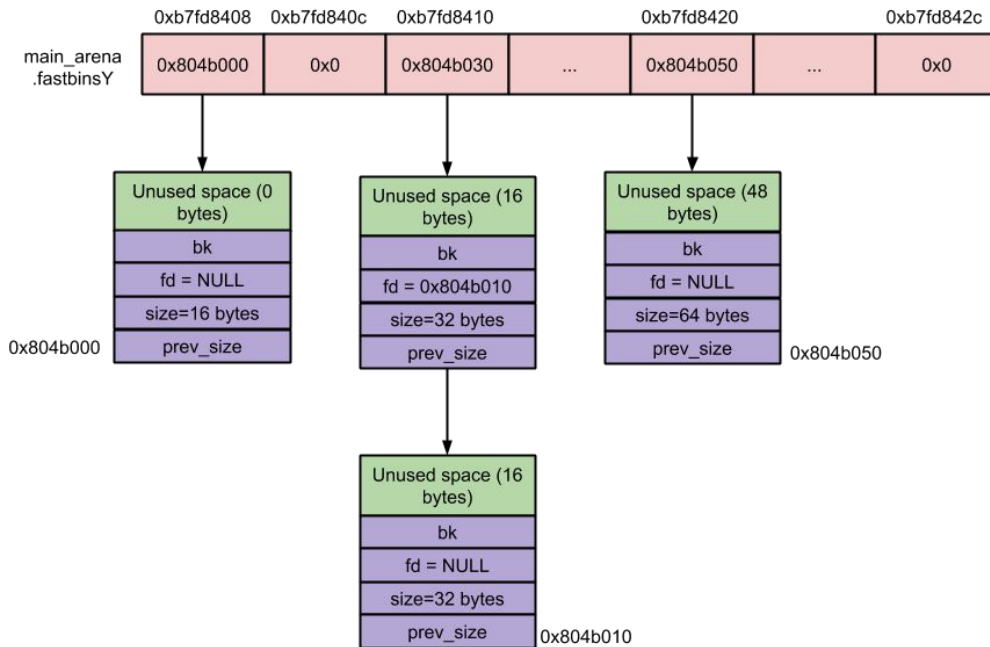
# Bins

- t-cache
- Fast bin (0x20 to 0x90 bytes)
- Unsorted bin
- Small bin (< 512 bytes )
- Large bin (>= 512 bytes)
- top-chunk

# Fast Bins

- Optimized bins for tiny freed chunks
- Managed as LIFO linked lists (backward pointer not used)
- Better performance, less checks and operations
- Freed chunks are never consolidate with any other freed chunk
- Freed chunks in fast bins act as non freed chunks (P flag of next chunk is not set to 0)





# Fast Bins



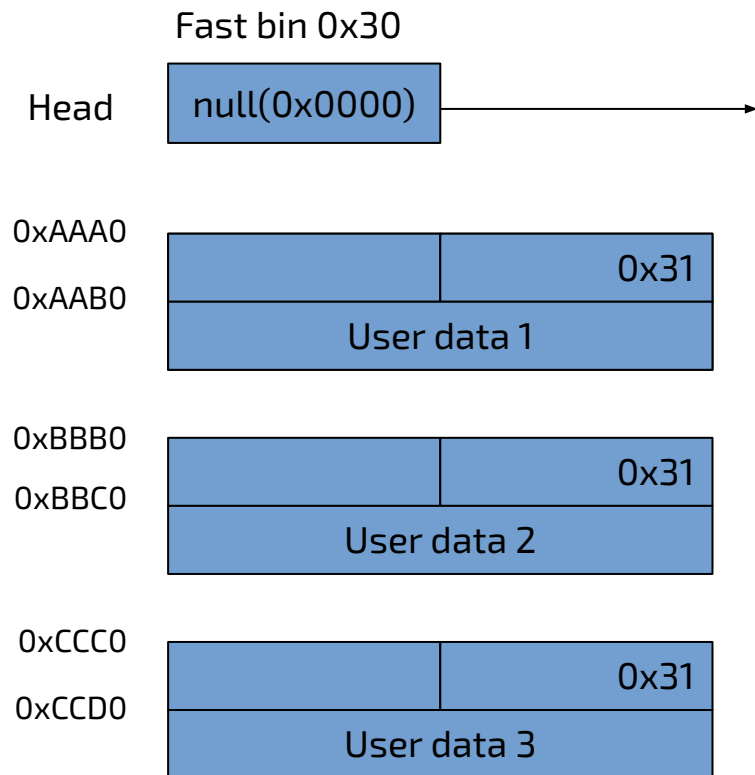
Fast Bin Snapshot

# Fast Bins

Fast bin 0x30




# Fast Bins: free(0xCCD0) - Before

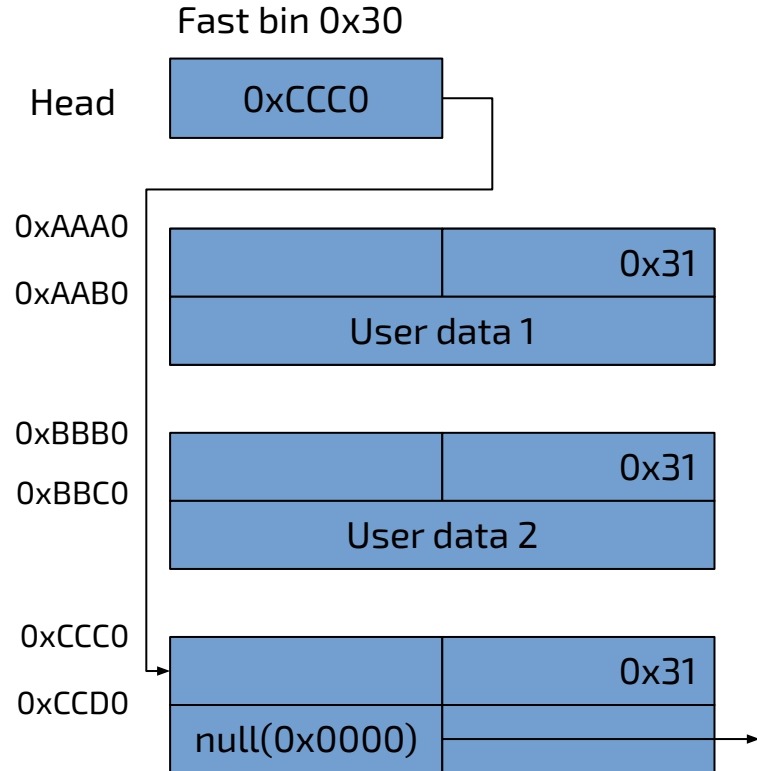





# Fast Bins: free(0xCCD0) - After



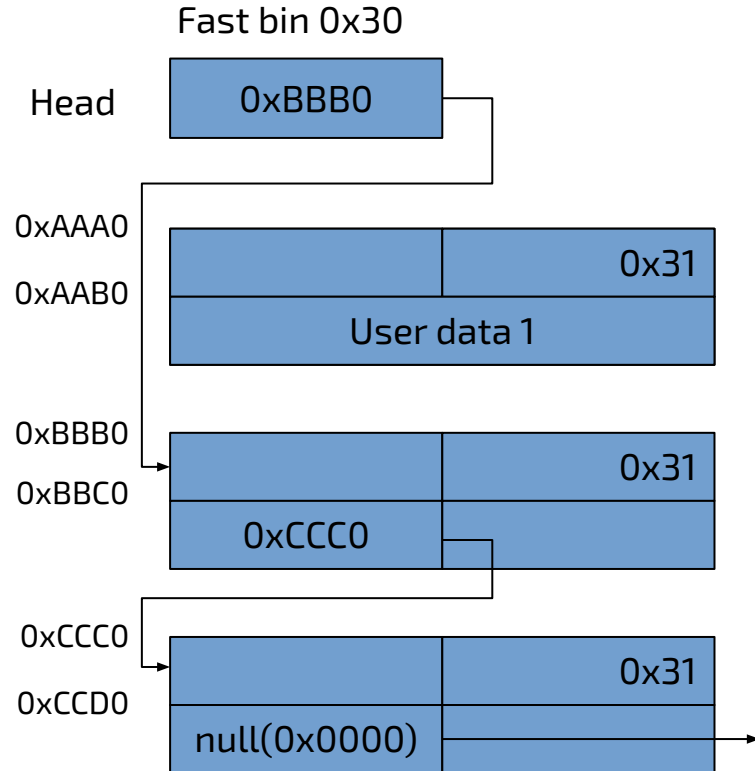
# Fast Bins: free(0xBBC0) - Before



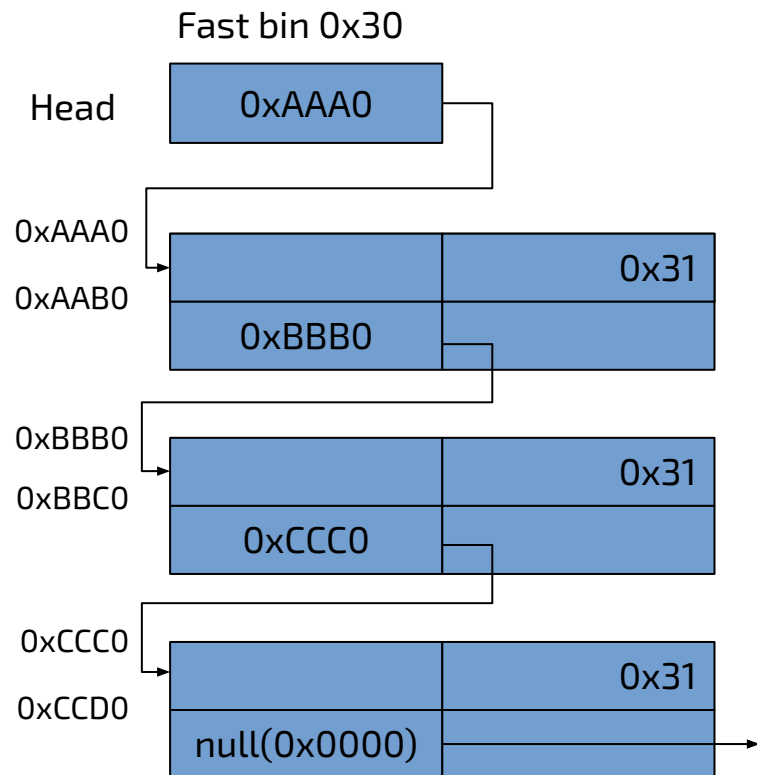
# Fast Bins: free(0xBBC0) - After



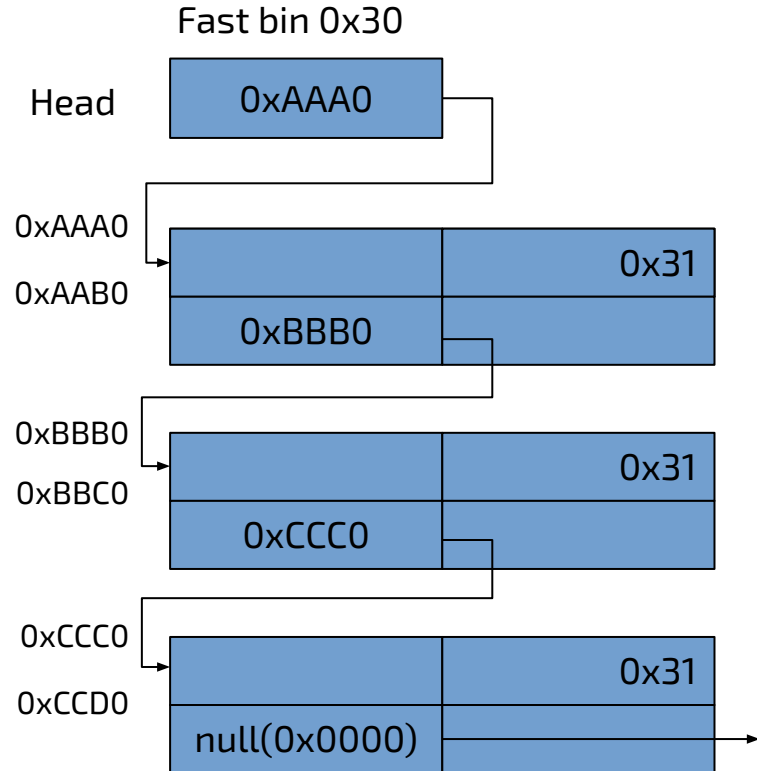
# Fast Bins: free(0xAAB0) - Before



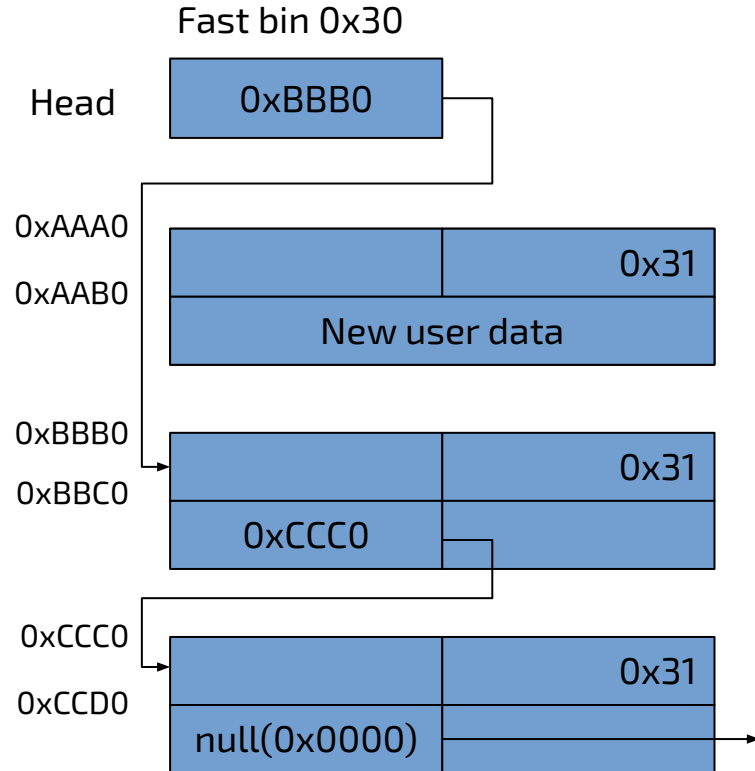
## Fast Bins: free(0xAAB0) - After



# Fast Bins: malloc(0x20) - Before



## Fast Bins: malloc(0x20) - After



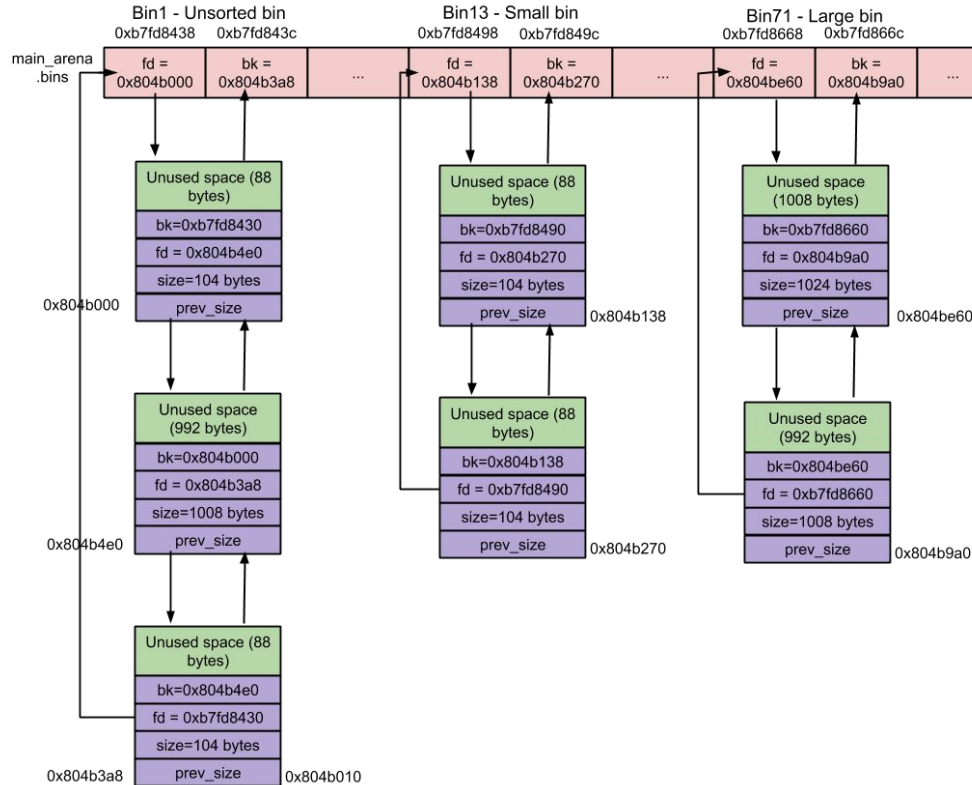
Let's see it in memory!



# Unsorted Bin

- Any freed chunk with size  $\geq 0xA0$  ends up in the unsorted bin
- Managed as a double linked list
- When a chunk in the unsorted bin is not able to satisfy a malloc request (e.g., `malloc(0x200)` but the freed chunk has size `0x100`), the chunk in the unsorted bin is moved to the proper small or large bin
- Unsorted bin is like a middle ground

# Bins (Unsorted, Small, Large )



Unsorted, Small and Large Bin Snapshot

Let's see it in memory!

# Vulnerability after Allocation

- Good old buffer overflow :D
- Overflows on:
  - Metadata and content of the next chunks ( in memory)
  - Top chunk (House of force)
- Potential leaks if the buffer is not memset to 0 (calloc solves this problem)

# Vulnerability after Deallocation




- Pointer should be set to 0 after free, otherwise it may occur:
  - Leakage of the bins' pointers
  - Corruption of the bins' pointers
  - Multiple pointers to the same chunk in memory
  - Double free (Fastbin attack)

# Fast Bin Attack




- Allocate twice the same chunk
- Allocate an almost arbitrary chunk

# Fast Bin Attack

Fast bin 0x30



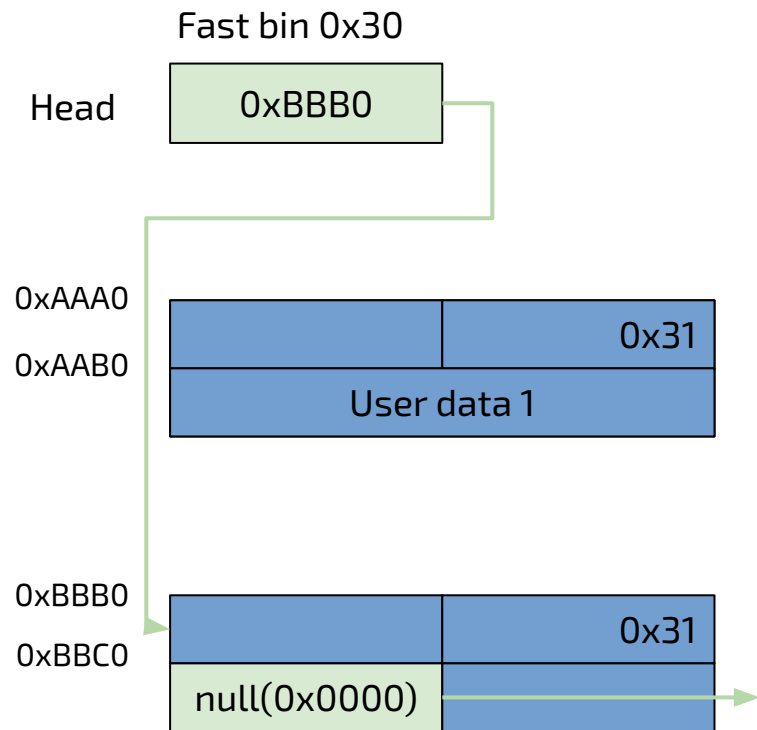
# Fast Bin Attack



`free(0xBBC0)`




# Fast Bin Attack




`free(0xBBC0)`

# Fast Bin Attack



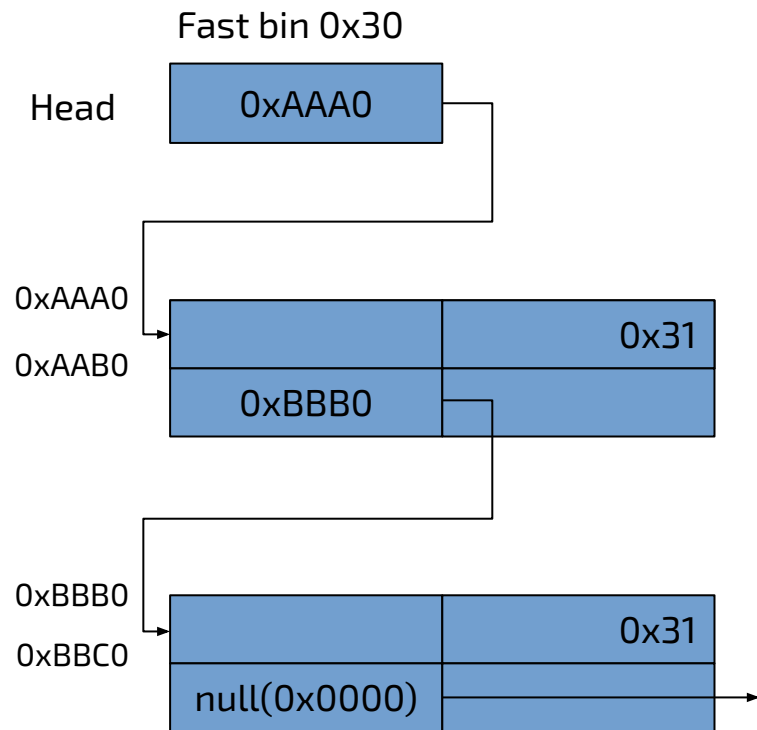
`free(0xAAB0)`

# Fast Bin Attack




`free(0xAAB0)`

# Fast Bin Attack



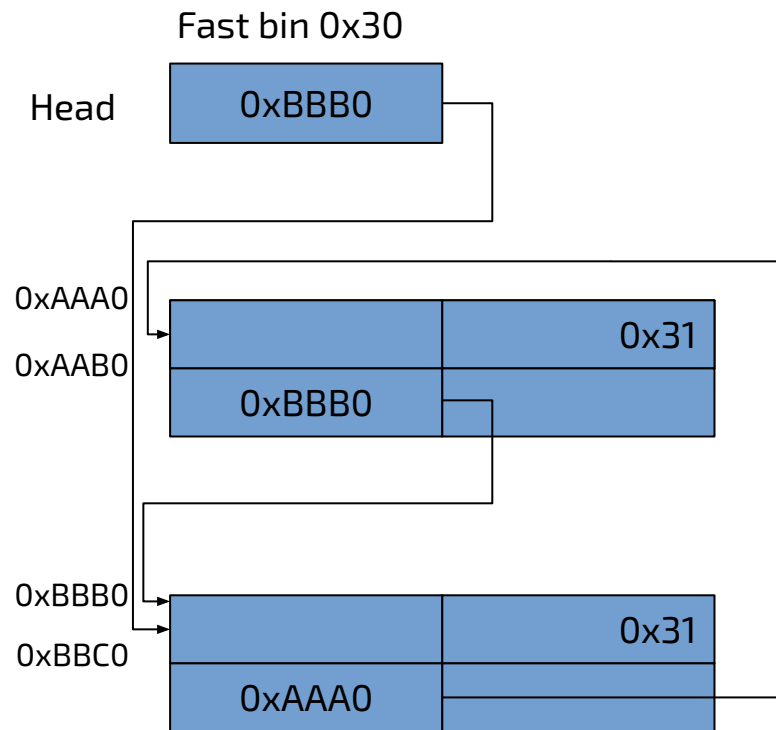
`free(0xBBC0)`

# Fast Bin Attack



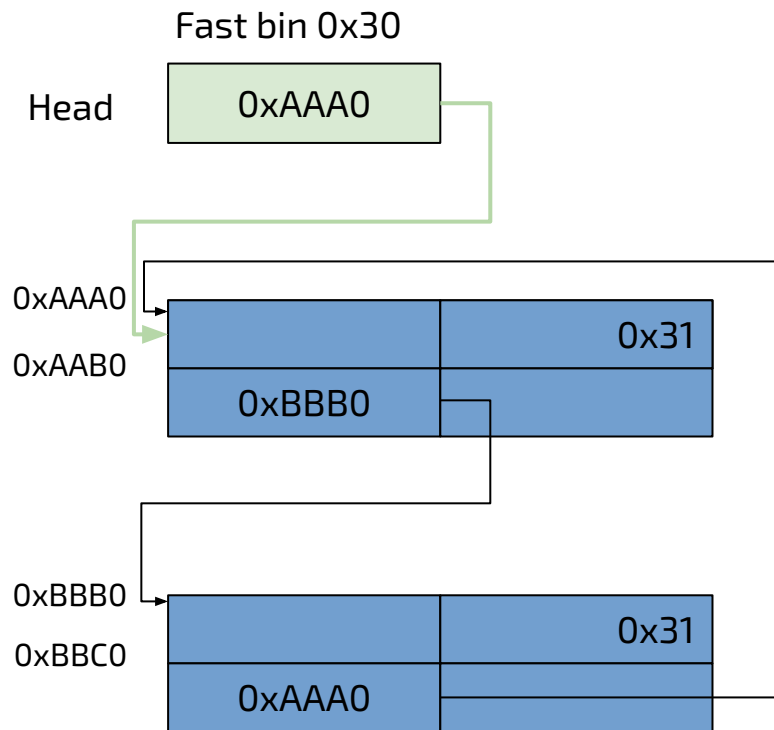
`free(0xBBC0)`

# Fast Bin Attack



`malloc(0x20)`

# Fast Bin Attack

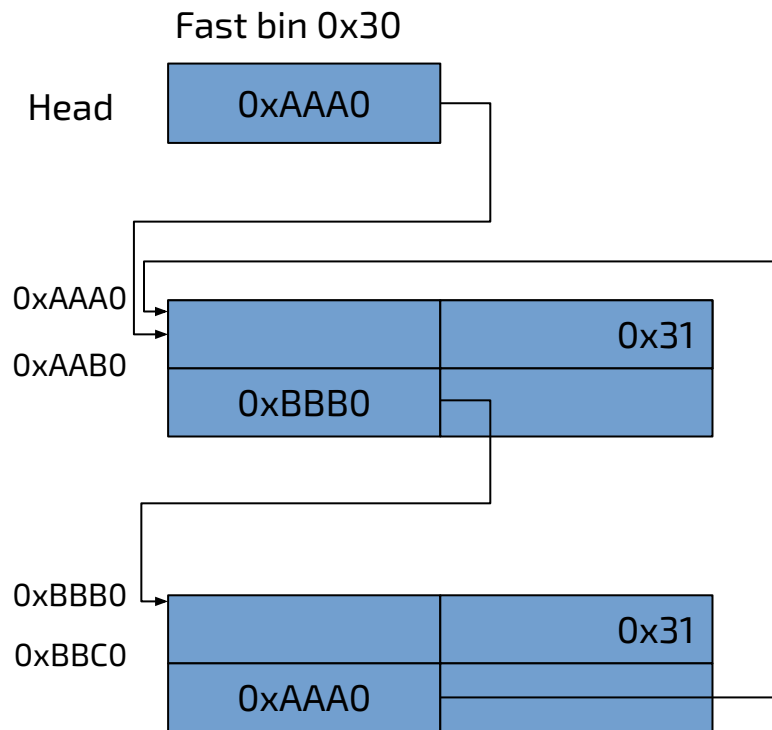


`malloc(0x20)`

0xBBC0 is returned by the malloc.

- Buffer 1: 0xBBC0

# Fast Bin Attack

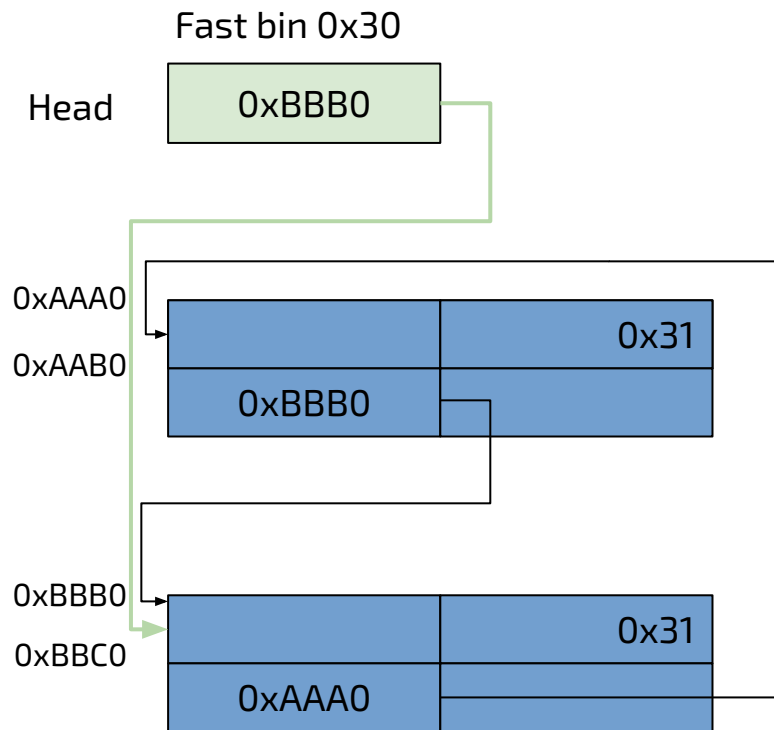


```
malloc(0x20)
```

- Buffer 1: 0xBBC0



# Fast Bin Attack

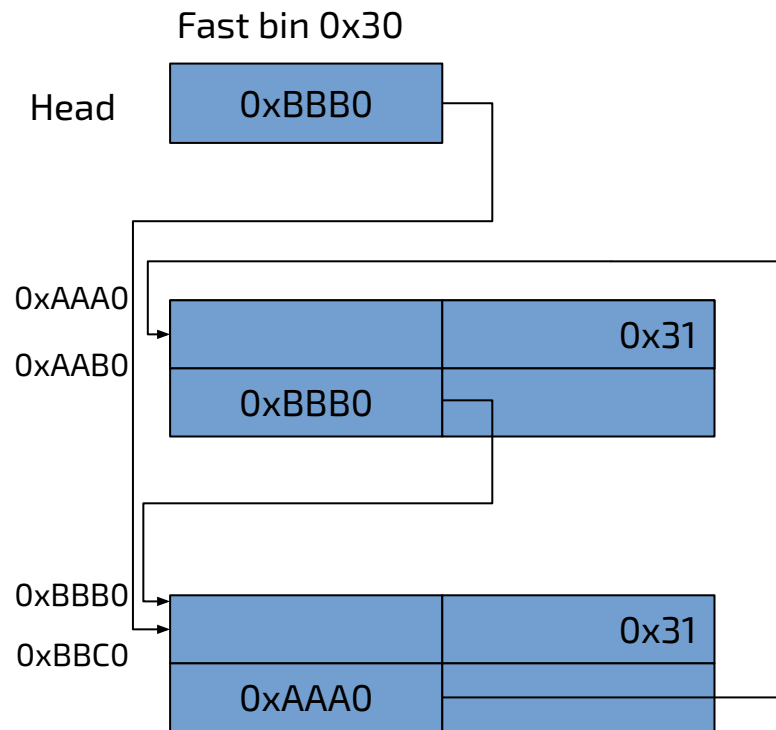


`malloc(0x20)`

0xAAB0 is returned by the malloc.

- Buffer 1: 0xBBC0
- Buffer 2: 0xAAB0


# Fast Bin Attack



write Buffer 1

- Buffer 1: 0xBBC0
- Buffer 2: 0xAAB0

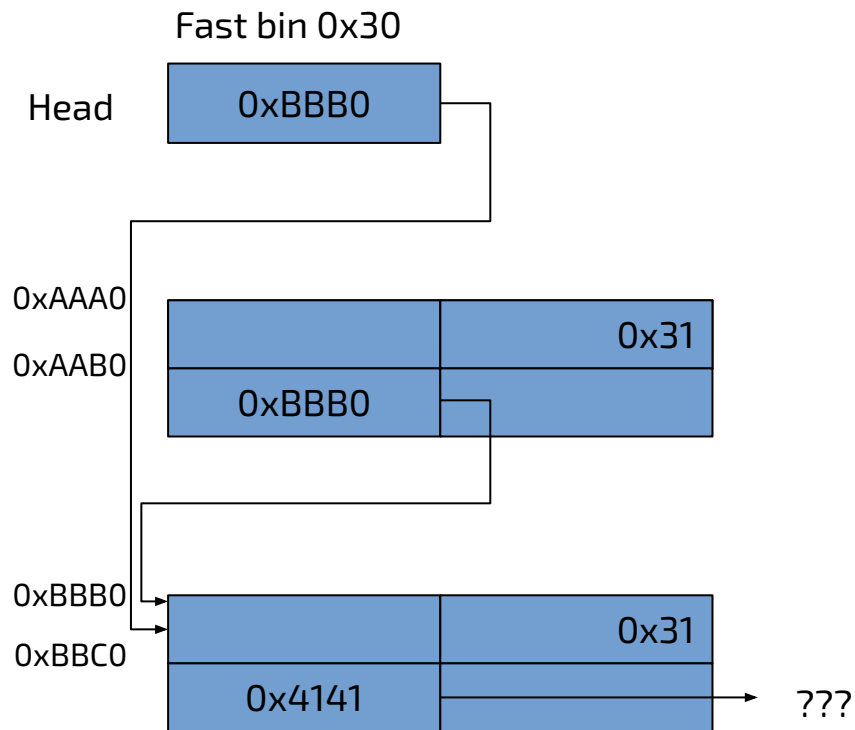
# Fast Bin Attack



write Buffer 1

- Buffer 1: 0xBBC0
- Buffer 2: 0xAAB0



# Fast Bin Attack



`malloc(0x20)`

- Buffer 1: 0xBBC0
- Buffer 2: 0xAAB0

# Fast Bin Attack: malloc(0x20) - After




## malloc(0x20)

0xBBC0 is returned by the malloc.

- Buffer 1: 0xBBB0
- Buffer 2: 0xAAB0
- Buffer 3: 0xBBC0

# Notes

- The arbitrary chunk will be returned if:
  - The memory is mapped (otherwise SEG FAULT)
  - The size of the fake chunk matches with the bin size
- The first 4 bits of the size are not considered
- No requirements on the alignment of the chunk



# Load another Library (libc-2.xx.so)

- **env LD\_PRELOAD**

- LD\_PRELOAD=./libc-2.23.so ./binary

- **ld.so**

- ./ld-2.23.so --library-path ./lib ./binary
- lib contains libc.so.6

- **patchelf** (<https://github.com/NixOS/patchelf>)

- patchelf --set-interpreter ./ld-2.23.so --replace-needed libc.so.6 ./libc-2.23.so ./binary

- **YOLO** (Do not use this!)

- Replace system library

- **Docker/Virtual Machine**

# Poison Null Byte


- Allocate two chunks that overlap




# Poison Null Byte

```
char *buf = malloc(128);  
int read_length = read(0, buf, 128);  
buf[read_length] = 0;
```


# Poison Null Byte




Initial Setup




Free( B )




Overflow into B. Sizes goes from 0x208 to 0x200.  
prev\_size is not update




Allocate two chunks into old B. First not a fastbin



Free( B1 )




Free( C ) trigger the merge with the previous chunk




the next allocation will overlap with B2

# Chunk



# Null Byte



# Poison Null Byte “FIX”

```
if ( __builtin_expect ( chunksize (P) != prev_size (next_chunk (P)), 0))  
    malloc_printerr ("corrupted size vs. prev_size");
```

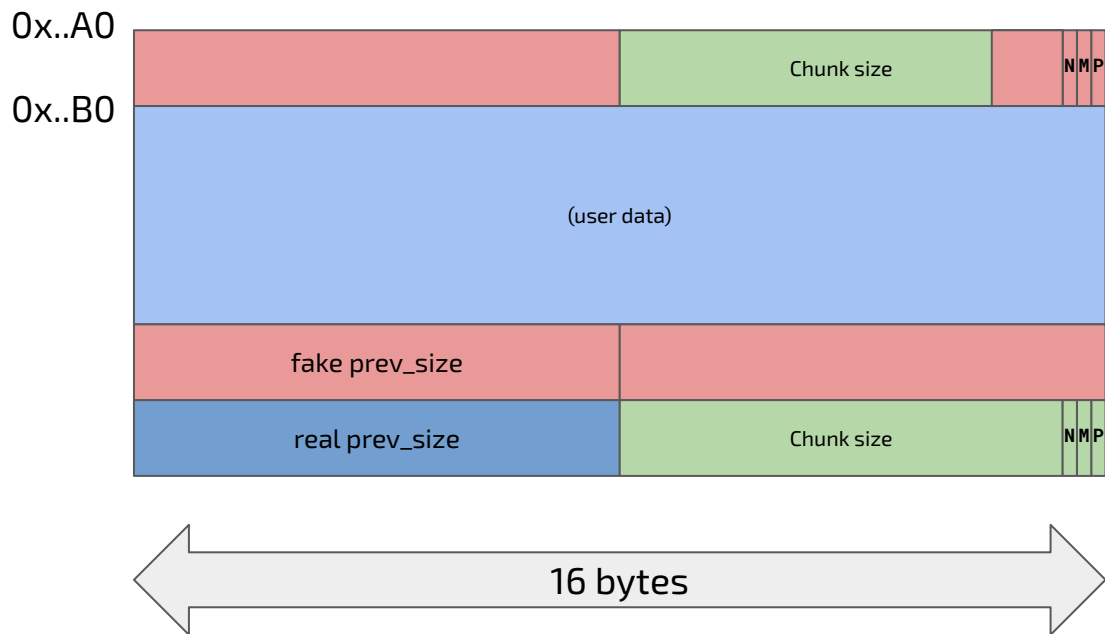
# Poison Null Byte “FIX”

```
if ( __builtin_expect (chunksize(P) != prev_size (next_chunk(P)), 0))  
    malloc_printerr ("corrupted size vs. prev_size");
```


```
/* Size of the chunk below P. Only valid if !prev_inuse (P). */  
#define prev_size(p) ((p)->mchunk_prev_size)
```

```
/* Ptr to next physical malloc_chunk. */  
#define next_chunk(p) ((mchunkptr) (((char *) (p)) + chunksize (p)))
```


# fake prev\_size




# Poison Null Byte Fixed




Initial Setup (B chunk needs to contain fake prev\_size)



Free( B )




Overflow into B. Sizes goes from 0x208 to 0x200.  
prev\_size is not update




Allocate two chunks into old B. First not a fastbin



Free( B1 )



Free( C ) trigger the merge with the previous chunk



the next allocation will overlap with B2





T-Cache

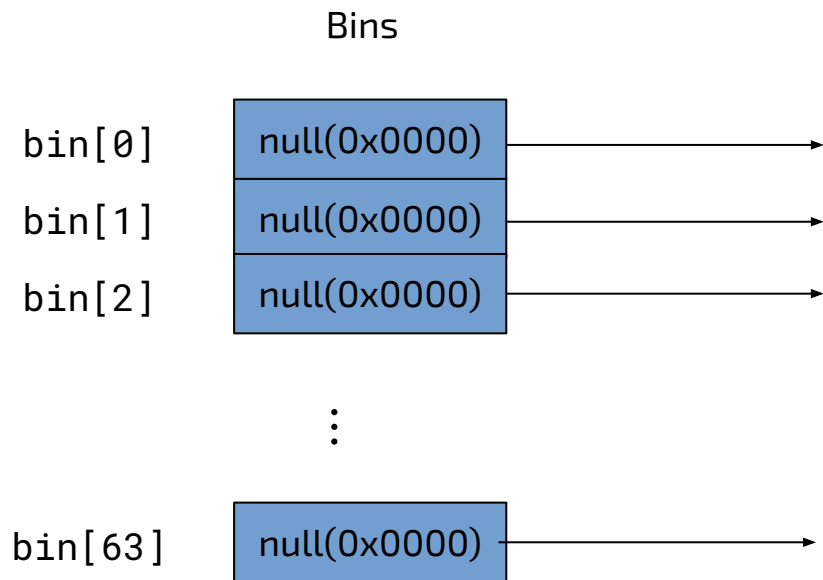
# T-Cache

- **Cache** for chunk size < **0x500**
- **LIFO**
- **New Attack Vector**
- Need to **bypass** it for attacks on other bins
- You can exploit T-Cache for better **HEAP Manipulation**
- **64 bins**
- **7** chunks per bin as cache

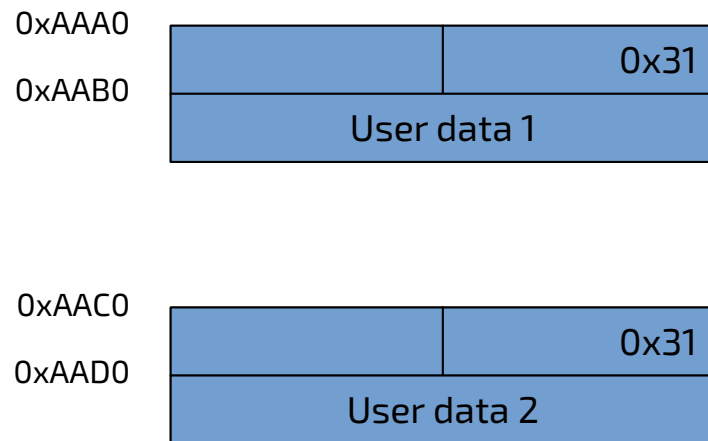
# T-Cache



# T-Cache




`free(0xAAB0)`




# T-Cache

`free(0xAAB0)`




# T-Cache

`free(0xAAC0)`



# T-Cache

`free(0xAAC0)`

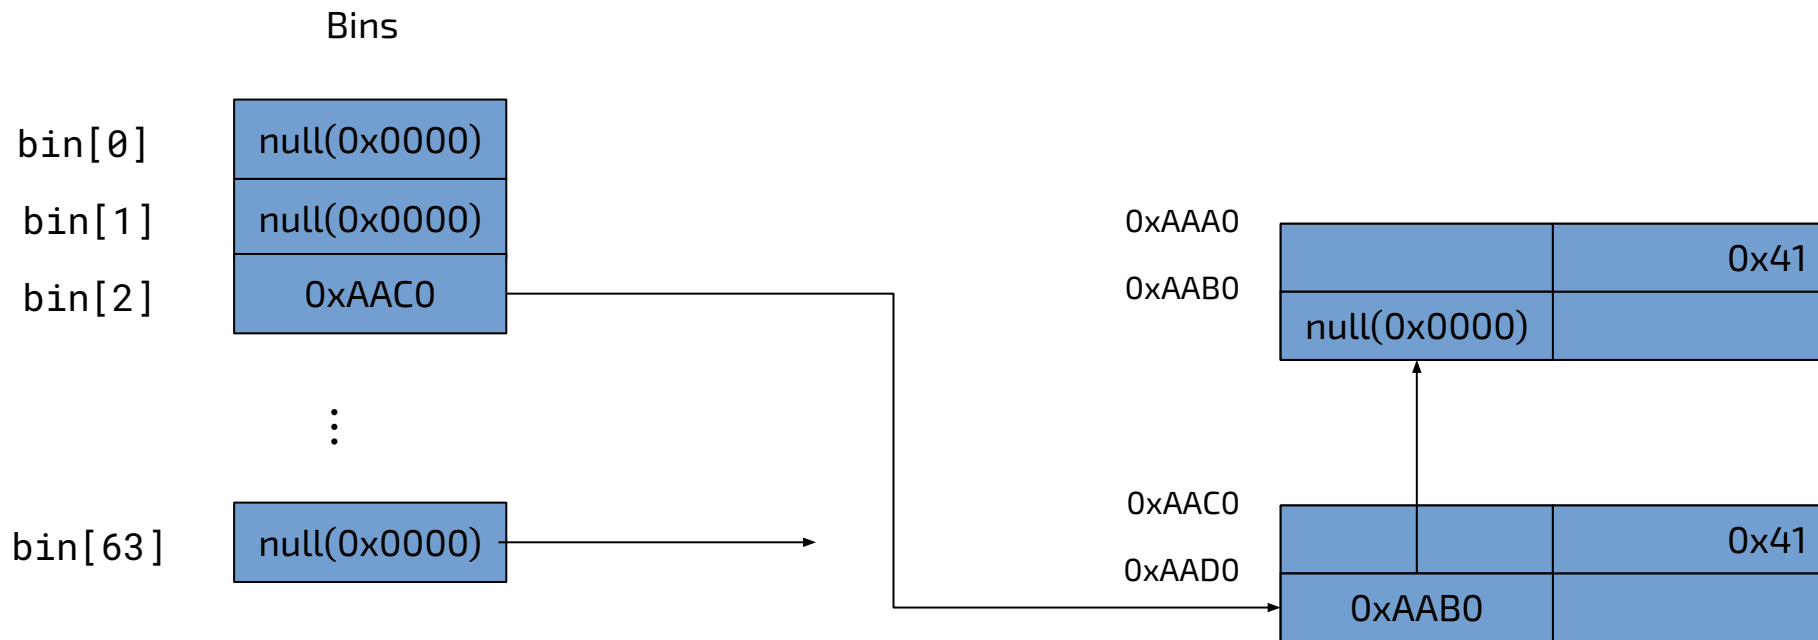


# T-Cache Poison


- Allocate an arbitrary chunk



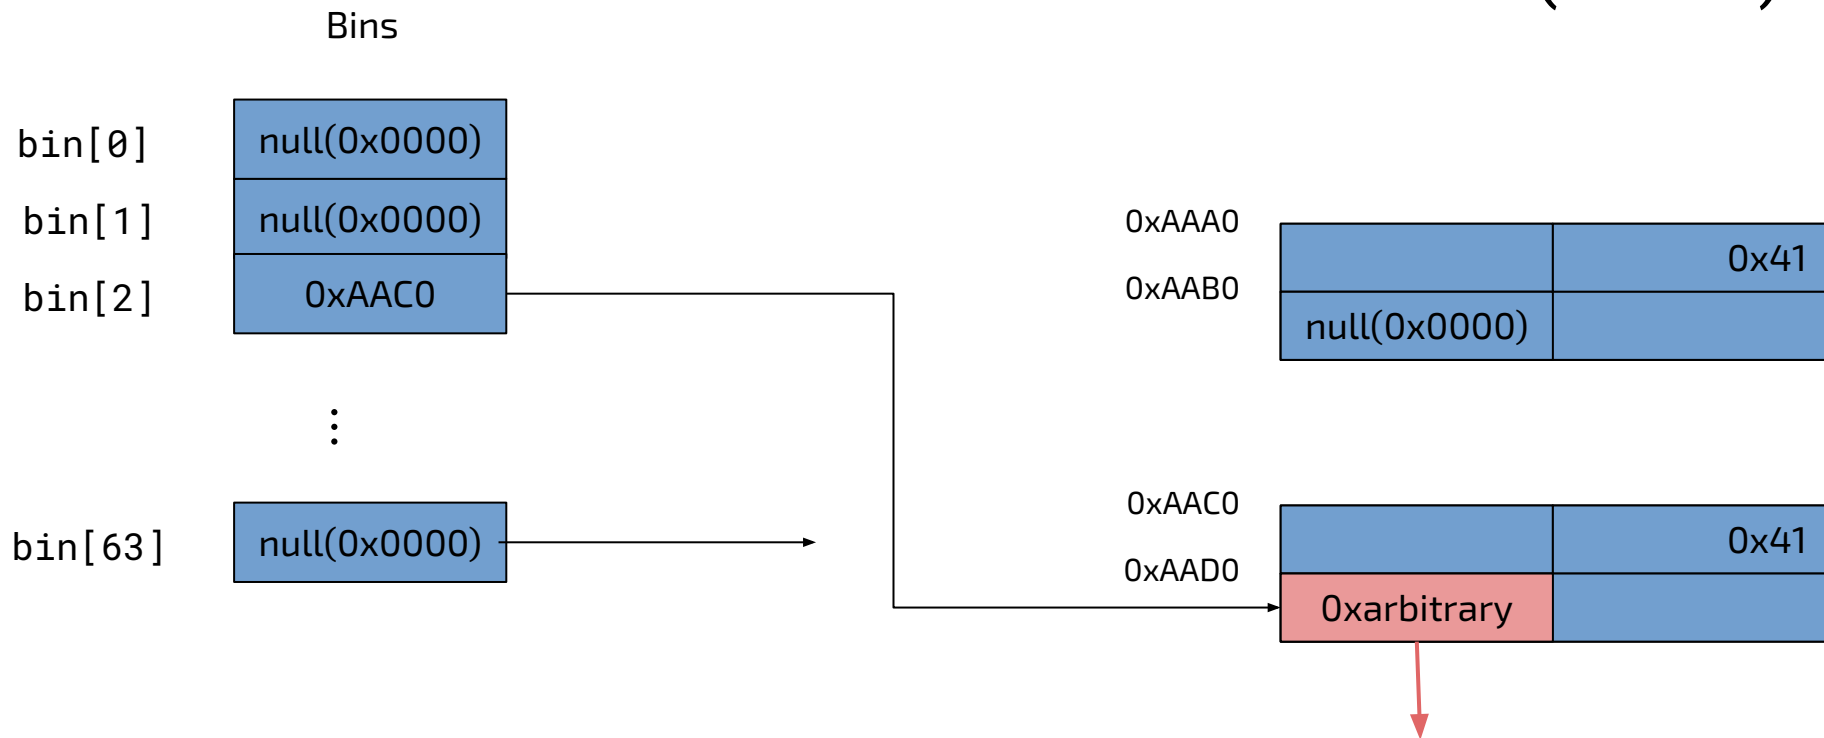
# T-Cache (need at least 2 elements in the list)



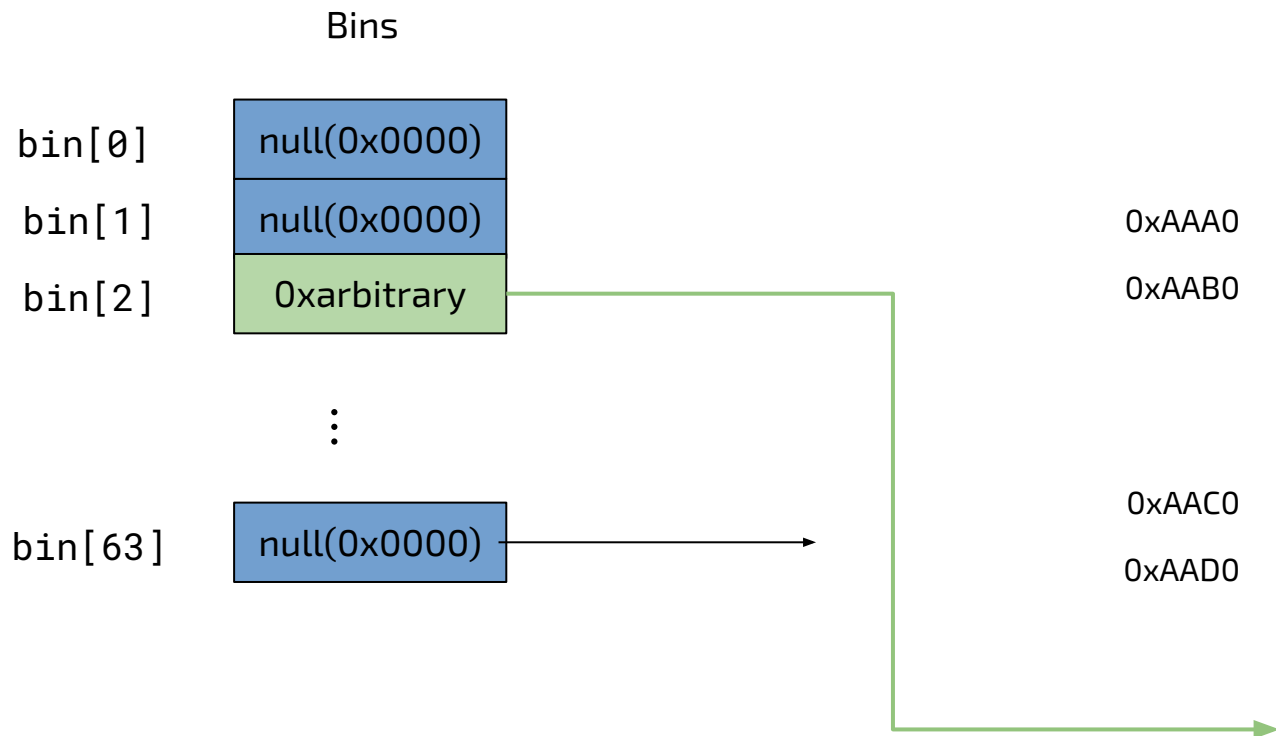
# T-Cache



# T-Cache




# T-Cache




malloc(0x20)

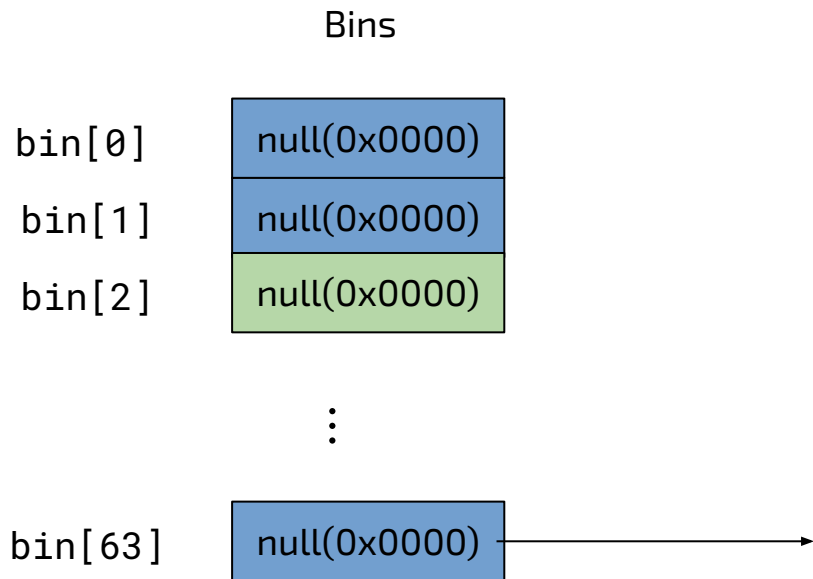
Chunk1 : 0xAAD0



# T-Cache





# T-Cache



## malloc(0x20)

Chunk1 : 0xAAD0

Chunk2 : **0xarbitrary**



# T-Cache Key > glibc 2.29

## glibc 2.34

```
tcache_key = random_bits ();  
tcache_key = (tcache_key << 32) | random_bits ();  
e->key = tcache_key;
```

## glibc 2.29 - 2.33

```
e->key = tcache;
```

## free

```
if (__glibc_unlikely (e->key == tcache)) {  
    /* very likely a problem make extra checks*/  
}
```

# T-Cache PTR Protection

glibc >= 2.32

```
e->next = PROTECT_PTR (&e->next, tcache->entries[tc_idx]);
```

```
#define PROTECT_PTR(pos, ptr) \  
    ((__typeof (ptr)) (((size_t) pos) >> 12) ^ ((size_t) ptr)))
```



# Best documentation is source code.


[illegible]

Best documentation is ~~source code~~. binary code

Ubuntu Backports - glibc 2.27

```
@@ -2942,6 +2951,7 @@ tcache_get (size_t tc_idx)
    assert (tcache->entries[tc_idx] > 0);
    tcache->entries[tc_idx] = e->next;
    --(tcache->counts[tc_idx]);
+   e->key = NULL;
    return (void *) e;
}
```

# Algorithm



<https://raw.githubusercontent.com/cloudburst/libheap/master/heap.png>

## Useful Links / Reading Material

- <https://github.com/shellphish/how2heap>
- <https://sploitfun.wordpress.com/2015/02/10/understanding-glibc-malloc/>
- <https://heap-exploitation.dhavalkapil.com>
- <https://www.usenix.org/conference/usenixsecurity18/presentation/heelan> (Automatic Heap Manipulation)